

Image Segmentation

- Introduction
- Strategies for image segmentation
- examples

Introduction

- Objective
 - to divide an image in regions or objects according to some criterium
 - groups of pixels sharing similar characteristics
 - it is often a crucial step in enabling the interpretation of an image
 - the segmentation accuracy often determines the success of the following interpretation steps
- Two generic strategies:
 - Discontinuity or edge-based
 - the image is partitioned based on sudden changes or discontinuities on the image characteristics
 - Similarity or region-based (spatial grouping/clustering of pixels)
 - partitioning occurs based on the similarity degree between neighbouring pixels

Discontinuity-based segmentation

- Strategy based on the assumption that between two regions there must be a boundary, frontier, discontinuity or edge
 - the idea is to identify those discontinuities and thus accurately delineate regions
- Edges thus represent significant local changes in the image intensity values
 - they are important features to analyse an image and enable separating regions, objects or to identify changes in illumination or colours
- Implemented using Local Edge Operators
 - Gradient of image intensity
 - Robert, Sobel and Prewitt operators
 - Second Derivative Operators
 - Laplacian, Laplacian of Gaussian

Similarity-based segmentation

- Region-based segmentation methods that attempt to partition or group regions according to common image properties
 - intensity values from original images
 - computed values based on an image operator
 - textures or patterns that are unique to each region
- Complex systems may use a combination of these properties to segment images
 - basically the idea is to group/cluster pixels that are similar or that have similar statistical properties
 - and/or separate pixels that are dissimilar

Discontinuity detection - image derivatives

- the concept is to compute the first or second derivatives of images
- the derivatives of an image are functions of an independent variable (coordinates/space) with infinitesimal increments

- the 1st derivative or gradient:

$$\nabla f(x,y) \equiv \delta f / \delta x + \delta f / \delta y = f'(x) + f'(y) = f(x+1) - f(x) + f(y+1) - f(y)$$

- if there are no changes in the intensity of the image then the derivative/gradient is zero
- before any change or immediately after a change its value is zero
- it is only during a change of intensity that its value is non-zero
- the 2nd derivative has non-zero value only at the instants where there is a transition from no change to change in the intensity of the image and vice-versa

Discontinuity detection - gradient

- The gradient is the first derivative of the image $f(x, y)$
 - it provides a measure of how the function $f(x, y)$ changes as a function of changes in the arguments x and y

$$\nabla f(x, y) \equiv \nabla_x + \nabla_y = G_x(x, y) + G_y(x, y) = \delta f / \delta x + \delta f / \delta y =$$

$$f'(x) + f'(y) = f(x+1) - f(x) + f(y+1) - f(y)$$

- the gradient vector points in the direction of maximum change
- its length indicates the magnitude of the gradient
- in practical terms, the computation of the derivatives can be approximated to the convolution of specific kernels/masks

first derivative / gradient

$$G_x = \delta f / \delta x = [-1, 1]$$

$$G_y = \delta f / \delta y = [-1, 1]^T$$

second derivative

$$\delta^2 f / \delta x^2 = [1, -2, 1]$$

$$\delta^2 f / \delta y^2 = [1, -2, 1]^T$$

Discontinuity detection - gradient

- The gradient is the first derivative of the image $f(x, y)$
 - it provides a measure of how the function $f(x, y)$ changes as a function of changes in the arguments x and y

$$\nabla f(x, y) \equiv \nabla_x + \nabla_y = G_x(x, y) + G_y(x, y) = \delta f / \delta x + \delta f / \delta y =$$

$$f'(x) + f'(y) = f(x+1) - f(x) + f(y+1) - f(y)$$

- the gradient vector points in the direction of maximum change
 - its length indicates the magnitude of the gradient
- So the operations to perform using kernels $H(x, y)$:
 - obtain horizontal edges by computing $G_x(x, y) = H_x * f(x, y)$
 - convolution with the horizontal filter kernel H_x
 - obtain vertical edges by convolving with the vertical kernel H_y

Discontinuity detection - gradient using local operators

- but normally, larger kernels are used to include smoothness and deliver coherent results in both directions
 - common operators for computing the gradient are the Roberts, Prewitt and Sobel kernels

Roberts	$H_x(x,y)$	$H_y(x,y)$	Prewitt	$H_x(x,y)$	$H_y(x,y)$																								
	<table><tr><td>0</td><td>1</td></tr><tr><td>-1</td><td>0</td></tr></table>	0		1	-1	0	<table><tr><td>1</td><td>0</td></tr><tr><td>0</td><td>-1</td></tr></table>	1	0	0	-1	<table><tr><td>-1</td><td>0</td><td>1</td></tr><tr><td>-1</td><td>0</td><td>1</td></tr><tr><td>-1</td><td>0</td><td>1</td></tr></table>	-1	0	1	-1	0	1	-1	0	1	<table><tr><td>1</td><td>1</td><td>1</td></tr><tr><td>0</td><td>0</td><td>0</td></tr><tr><td>-1</td><td>-1</td><td>-1</td></tr></table>	1	1	1	0	0	0	-1
0	1																												
-1	0																												
1	0																												
0	-1																												
-1	0	1																											
-1	0	1																											
-1	0	1																											
1	1	1																											
0	0	0																											
-1	-1	-1																											
Sobel			$H_x(x,y)$	$H_y(x,y)$																									
			<table><tr><td>-1</td><td>0</td><td>1</td></tr><tr><td>-2</td><td>0</td><td>2</td></tr><tr><td>-1</td><td>0</td><td>1</td></tr></table>	-1	0	1	-2	0	2	-1	0	1	<table><tr><td>1</td><td>2</td><td>1</td></tr><tr><td>0</td><td>0</td><td>0</td></tr><tr><td>-1</td><td>-2</td><td>-1</td></tr></table>	1	2	1	0	0	0	-1	-2	-1							
-1	0	1																											
-2	0	2																											
-1	0	1																											
1	2	1																											
0	0	0																											
-1	-2	-1																											

- So the operations to perform using the kernels $H(x,y)$ to compute the gradient are:
 - obtain horizontal edges by computing $G_x(x,y)=H_x*f(x,y)$
 - this is the convolution with the horizontal filter kernel H_x
 - obtain vertical edges by convolving with the vertical kernel H_y

Discontinuity detection - gradient using local operators (2)

- Using the gradient or derivatives, two values are computed

- magnitude $|\nabla f(x,y)| \equiv G$ and orientation θ

$$G = \sqrt{(G_x^2 + G_y^2)} \quad \theta = \arctan(G_y/G_x)$$


or simply $G = G_x + G_y$ or $\max(G_x, G_y)$

- if magnitude is above a certain threshold, the pixel is considered to be an edge
- but a further decision needs to be made
 - whether two pixels that have been identified as edges belong to the same contour/frontier
- Often, the different pixels that have been identified as edges, do not define a unified boundary/frontier due to
 - noise in the picture, breaks in the boundary caused by occlusion, etc
 - consequently the edge detection operation is followed by a linking operation to assemble edge pixels into uniform boundaries

Discontinuity detection - edge linking

- 2 alternative strategies can be used
 - Local, comparing magnitude and orientation of the gradient of both pixels
 - Global, using the Hough transform

Discontinuity detection - edge linking (2)



Discontinuity detection - second derivative

- The Laplacian operator highlights regions in an image of rapid intensity change by calculating the second spatial derivative

$$L(x,y) = \delta^2 f / \delta x^2 + \delta^2 f / \delta y^2$$

- the operator is normally applied to a gray scale image producing another gray scale image as output
- The calculation of the second derivatives of an image can be approximate by standard convolutional filters represented by their kernels/masks as for the gradient case
- Two commonly used small kernels using a negative peak are

0	-1	0
-1	4	-1
0	-1	0

-1	-1	-1
-1	8	-1
-1	-1	-1

- Using one of these kernels, the Laplacian can be calculated using standard convolution methods

Discontinuity detection - second derivative (2)

- kernels that approximate a second derivative measurement on the image are very sensitive to noise (they detect “instantaneous” transitions)
- as such the image is often filtered (Gaussian filter) to reduce high frequency noise, prior to applying the Laplacian masks
- given that spatial filtering can be implemented through convolution and because convolution is associative
 - the Gaussian smoothing filter is first convolved with the Laplacian filter and only then this hybrid filter is convolved with the image
 - the Laplacian of Gaussian approach, LoG

Discontinuity detection - LoG operator

- Given that both the Gaussian and the Laplacian kernels are usually much smaller than the image, this method leads to fewer arithmetic operations
- The LoG kernel can be calculated beforehand
 - thus, only one convolution needs to be performed at run-time on the image
- the usual 2-D LoG function is centred on zero with Gaussian standard deviation σ has the form:

$$LoG(x, y) = -\frac{1}{\pi\sigma^4} \left[1 - \frac{x^2 + y^2}{2\sigma^2} \right] e^{-\frac{x^2 + y^2}{2\sigma^2}}$$

Discontinuity detection - LoG operator (2)

- discrete kernel that approximates this function (using a Gaussian filter with $\sigma = 1.4$)

0	1	1	2	2	2	1	1	0
1	2	4	5	5	5	4	2	1
1	4	5	3	0	3	5	4	1
2	5	3	-12	-24	-12	3	5	2
2	5	0	-24	-40	-24	0	5	2
2	5	3	-12	-24	-12	3	5	2
1	4	5	3	0	3	5	4	1
1	2	4	5	5	5	4	2	1
0	1	1	2	2	2	1	1	0

- when the Gaussian is made increasingly narrower, the LoG kernel becomes the same as the simple Laplacian kernels shown previously
 - because smoothing a discrete grid/matrix with a very narrow Gaussian ($\sigma < 0.5$ pixels) has no effect
 - the simple Laplacian can be seen as a limiting case of the LoG for narrow Gaussians

Discontinuity detection - kernels for convolution

- Different kernels/masks may be used with different aims

-1	-1	-1
2	2	2
-1	-1	-1

horizontal

-1	-1	2
-1	2	-1
2	-1	-1

+45°

-1	2	-1
-1	2	-1
-1	2	-1

vertical

2	-1	-1
-1	2	-1
-1	-1	2

-45°

- the above represented masks are suitable for detecting bright lines above a dark background
- for dark lines above a bright background, symmetrical values should be used
- to detect both types, then the absolute values of the different convolutions' results should be used

-2	-1	0
-1	0	1
0	1	2

+45°

0	1	2
-1	0	1
-2	-1	0


-45°

Discontinuity detection - kernels for convolution (2)

- line detection example:



diagonal mask




Laplacian filter

0	-1	0
-1	4	-1
0	-1	0

-1	-1	-1
-1	8	-1
-1	-1	-1



Discontinuity detection - kernels for convolution (3)




Discontinuity detection - kernels for convolution (4)

- detecting isolated points


$$f(x-1,y-1) + f(x-1,y) + f(x-1,y+1) + f(x,y-1) - 8f(x,y) + f(x,y+1) + f(x+1,y-1) + f(x+1,y) + f(x+1,y+1)$$

1	1	1
1	-8	1
1	1	1



Discontinuity detection - challenges

- frontiers between objects in real-world images are usually not abrupt but rather gradual in terms of brightness
 - may be noisy or naturally smooth



Discontinuity detection - challenges (2)

- 4 different types of edges are normally considered



step

the intensity
change is
instantaneous
(ideal)



ramp

the intensity change is
not instantaneous but
rather occurs over a
finite distance



ridge / line

the intensity change is
instantaneous and
within a short
distance it returns to
its previous value

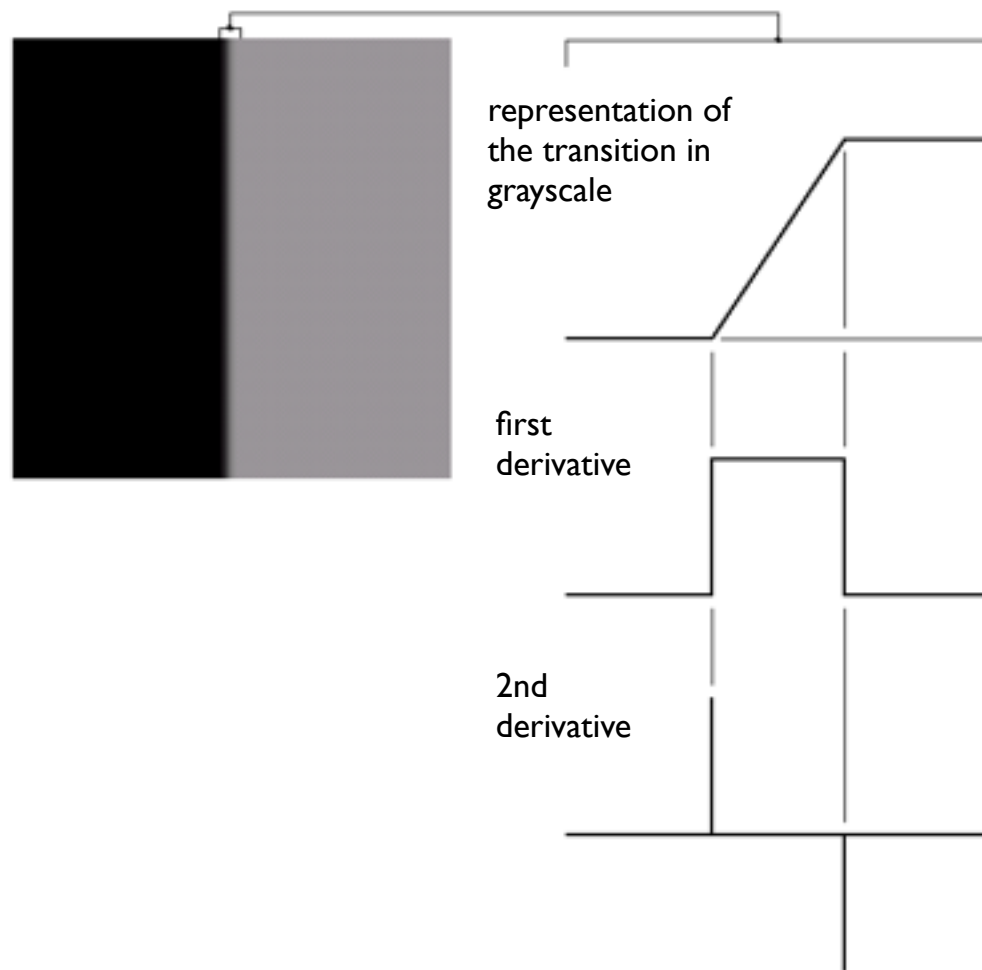


roof

the intensity change
occurs over a finite
distance and then it
returns back to its
previous value over a
finite distance

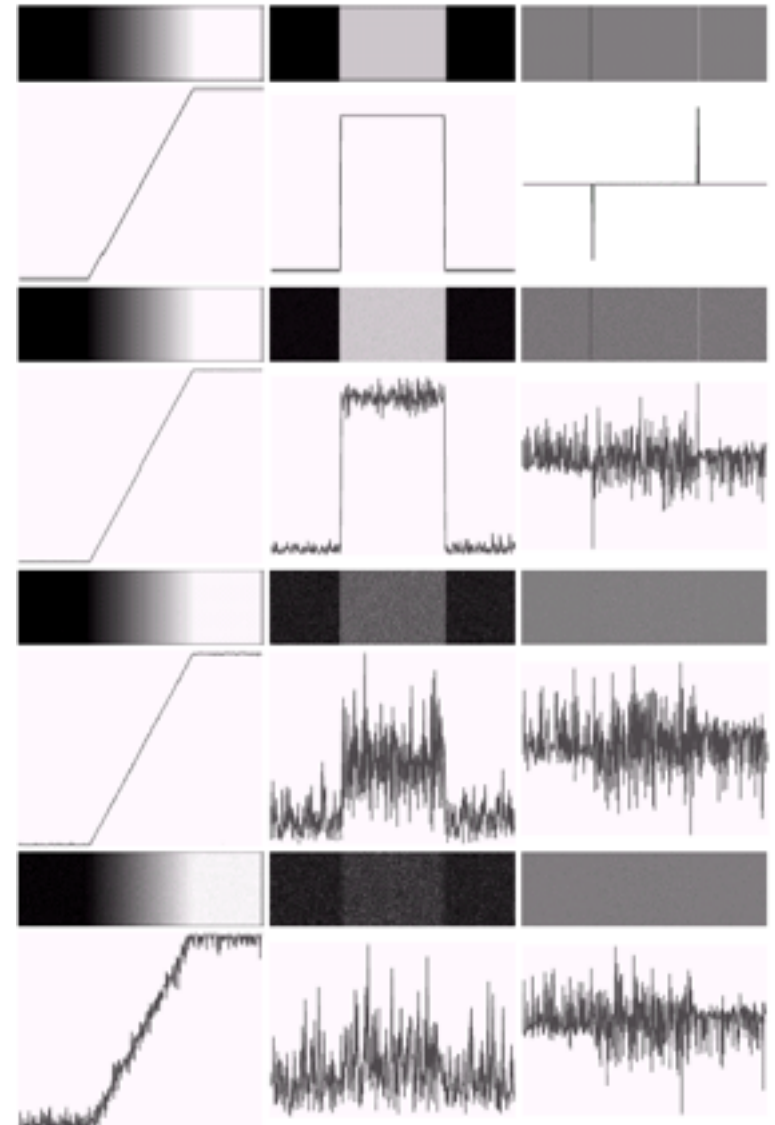
Discontinuity detection - challenges (3)

- it is easier to calculate the derivatives using a grayscale representation



discontinuity detection - challenges (4)

- the values of the derivatives may become too noisy to enable distinguishing an edge
- thus not allowing a decision to be taken
- the solution passes by performing a pre low-pass filtering to eliminate noise in the image
- needs to be performed in such a way as not to eliminate real edges



Detection by similarity (*clustering*)

- base concept is to identify regions in an image, to partition or group regions according to common image properties
- a region is defined as
 - an homogeneous part of the image in relation to some characteristic such as
 - intensity levels in a try scale
 - color values
 - texture
- basically the idea is to group/cluster pixels that are similar or that have similar statistical properties
 - and/or separate pixels that are dissimilar

Detection by similarity (*clustering*) (2)

- the inside of a region must be simple without many “holes” even if small in size (“hole” \equiv small aggregations of pixels that do not share the same homogeneous characteristic as the other pixels in the region)
- regions that are labelled as disjoint in a segmentation by similarity process, must present significantly distinct in relation to the characteristic in which one of them is homogeneous

Detection by similarity - binarization

- An import step in this segmentation approach is the binarization of the image
 - identification of pixels that are in the foreground and thus belong to objects from those that are in the background
 - it can be seen as a basic way of segmenting the image in only two groups
 - a black background and white objects in the foreground
 - binarization by threshold is one of the most common ways of achieving it
- the aim in binarization is to group pixels in only two classes, Q_1 and Q_2 (the background B and the foreground F)
 - with respect to some similar characteristic
 - necessary to find the value T that distinguishes those classes

Detection by similarity - binarization (2)

- there are three possible strategies to implement threshold binarization
 - global
 - uses only one threshold value for all pixels depending only on the overall intensity values
 - local
 - divides the image in several sub-images, applying a specific threshold value for each sub-image depending on their own characteristics
 - dynamic
 - the threshold value depends on the pixel position

Binarization by global thresholding

- regions/objects are separated using the pixel intensity
 - Global simple
 - one color is assigned to the background and another unique color to all the objects in the foreground

$$g(x,y) = \begin{cases} 0 & \text{se } f(x,y) < T \\ 1 & \text{se } f(x,y) \geq T \end{cases}$$

- Global multiple
 - one color is assigned to the background and different colours to each object in the foreground

$$g(x,y) = \begin{cases} 0 & \text{se } f(x,y) < T_1 \\ 1 & \text{se } T_1 \leq f(x,y) < T_2 \\ 2 & \text{se } T_2 \leq f(x,y) < T \\ \dots & \end{cases}$$

General approach to Global Binarization

1. estimate an initial global threshold T
2. segment the image by binarization using T
 - two groups are obtained: $G1$ with pixels $> T$ and $G2$ with pixels $\leq T$
3. calculate the average intensity for each group $m1$ and $m2$
4. calculate new $T = (m1 + m2)/2$
5. repeat 2 to 4, until the T between successive iterations is below a defined ΔT

Note 2: Matlab has an implementation for binarization in the function `imbinarization()`

General approach to Global Binarization (2)

- the threshold T is often estimated based on the average background intensity value
- which can be identified by several means
 - through the histogram
 - in principle the background is the largest region in the image so, it will correspond to the largest number of occurrences in the histogram
 - by subtracting two consecutive, or close in time, images in a video sequence
 - by the texture or color
 - etc

Note: a frequently used approach is the Otsu algorithm (next slides); Matlab has implementations for computing the threshold: `graythres()`, `otsuthresh()`, `adaptthresh()`

Automated approaches to estimate the threshold

- five ways of using the image histogram and identifying two distinctive modes, one for the background and one for the foreground

(I) image characteristics known - **cumulative histogram**

- if the objects in the foreground **occupy a fraction** I/p of the image and are **brighter** than the background

$$c(a) = \sum_{a=0}^A p(a) \quad T \text{ is chosen such that } c(T) = I/p$$

```
//input the value of p and the image dimensions NxM
for (i=0;i<M,i++)
for (j=0;j<N,j++)
    hist[I[i,j]]++;
// in MATLAB: imhist(I)
sum=0; a=0;
while(sum<(N*M)/p && a<256)
    { sum+=hist(a); a++; }
T = a-1; // desired threshold
```

- if the objects are **darker** than the background, **choose T** such that $c(T) = I - I/p$

Automated approaches to estimate the threshold (2)

(2) finding peaks and valleys in the histogram

- find the local maxima and then the minimum between them
 - because there may be noise, first the histogram is smoothed

(3) clustering by K-means variation (next slide)

Automated approaches to estimate the threshold (3)

(3) clustering by K-means variation

- starts by defining a threshold T and group pixels in 2 clusters
 - T is set to be equidistant to the averages of the intensities of the background and the foreground, respectively μ_B and μ_F
 - the background is assumed as being formed only by the four corner pixels, everything else being part of the foreground
- then a recursive cycle is initiated to minimize the distance between each pixel in a cluster and the average of that cluster
 - moving the pixel when necessary and recomputing the cluster intensity averages as well as the threshold
- so the basic idea is to cluster pixels according to how close their intensities are from the clusters' average intensities, μ_B and μ_F

$$\forall p \geq T : |p - \mu_B(T)| > |p - \mu_F(T)| \quad \&\& \quad \forall p < T : |p - \mu_B(T)| < |p - \mu_F(T)|$$

Automated approaches to estimate the threshold (4)

(4) clustering by the Otsu method

- the aim in binarization is to group pixels in only two classes, Q_1 and Q_2 obtaining the best value T that distinguishes those classes
- the Otsu algorithm selects T trying to minimize the variance inside each class (*intra-class*)

- probabilities $q_1(t) = \sum_{i=1}^t P(x)$ $q_2(t) = \sum_{i=t+1}^I P(x)$
- variances $\sigma_1^2(t) = \sum_{i=1}^t [i - \mu_1(t)]^2 \frac{P(i)}{q_1(t)}$ $\sigma_2^2(t) = \sum_{i=t+1}^I [i - \mu_2(t)]^2 \frac{P(i)}{q_2(t)}$
- averages $\mu_1(t) = \sum_{i=1}^t \frac{iP(i)}{q_1(t)}$ $\mu_2(t) = \sum_{i=t+1}^I \frac{iP(i)}{q_2(t)}$

- which is equivalent to maximising the inter-class variance

$$\sigma_b^2 = \sigma^2 - \sigma_w^2(t) = q_1(t)[1 - q_1(t)][\mu_1(t) - \mu_2(t)]^2$$

$$T: \max \sigma_b^2(t) = \sigma_b^2(T)$$

Automated approaches to estimate the threshold (4-2)

(4) clustering by the Otsu method (cont.)

- for each new computed value of T , $T: \max \sigma_b^2(t) = \sigma_b^2(T)$
 - separate the pixels into two clusters according to T
 - find the mean of each cluster
 - square the difference between the means
 - multiply by the number of pixels in one cluster times the number of pixels in the other (or the probabilities)
 - compute new threshold

$$n_B(T + 1) = n_B(T) + n_T \qquad \mu_B(T + 1) = \frac{\mu_B(T)n_B(T) + n_T T}{n_B(T + 1)}$$

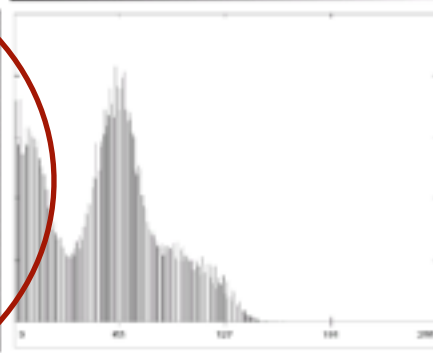

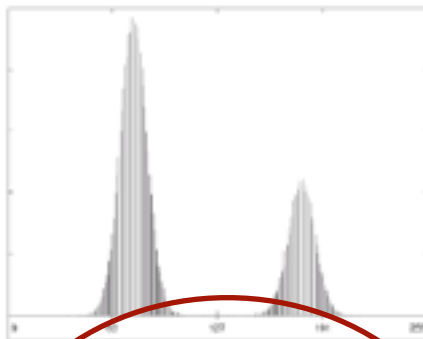

$$n_F(T + 1) = n_F(T) - n_T \qquad \mu_F(T + 1) = \frac{\mu_F(T)n_F(T) - n_T T}{n_F(T + 1)}$$

(5) Mixture Modelling

- assumes that the two classes - background and foreground - are Gaussian-distributed with certain mean and standard deviation, μ and σ , which are independent of the selected threshold
 - requires to compute an initial estimate of the Gaussian distributions (of the quantities n_B , n_F , μ_B , μ_F , σ_B , and σ_F)
 - the optimal threshold is the one that causes the mixture of the two estimated Gaussian distributions to best approximate the actual image histogram

Challenges in Global binarization

- local variations in light intensity (originated by shadows, etc)



- in this case, a global threshold is not adequate
- a possible solution is to pre-process the image to compensate such lightning variations
- equivalent to using a local thresholding technique in which the threshold value varies according to the position in the image

Challenges in binarization/thresholding

- only the individual pixel intensity is considered
 - not any relationship between pixels
 - no guarantee that the pixels identified by the thresholding process are contiguous (**connected/adjacent**)
- extraneous pixels are often included in regions to which they do not belong
 - on the other hand, especially near the boundaries of a region, isolated pixels that belong to the region may be missed
- in the presence of noise these undesired effects are aggravated
 - because there will be a higher probability that the pixel intensity is wrong (noisy), moving away from the normal intensity in the region
- lightning conditions/shadows also aggravate these errors

Challenges in binarization/thresholding (2)

- when using thresholding it is necessary to estimate the value of the threshold to distinguish between classes
- it may be necessary to “play” with different values
 - due to noise, lightning conditions but also due to the actual pixel intensities and real similarity of values between pixels in the foreground and background
 - some values may lead to excessive loss of pixels in a region, others to the inclusion of extraneous background pixels in the object/foreground

Region-based segmentation approaches

- region growing
 - the process starts with the selection of pixel “seeds”, to which more pixels are added when certain established criteria is observed
 - the criteria is specific to each problem
 - it can be (again) the simplest property that pixels in a region share - the intensity or the color
 - or higher order statistics of the region
 - necessary to define a stop criterium
 - necessary a careful selection of “seeds” as it may influence the final result as well as the processing delay

Region-based segmentation approaches (2)

- region growing
 - it may start considering the image as a single region, thus with only one seed
 - in which case it is normally the upper left corner
 - or with a set of regions represented by a set of seed points
 - analyses the values of neighbour pixels
 - those considered connected/adjacent are added to the region, making the region to grow
 - if a pixel is not connected it becomes the seed of a new region and the analysis process re-starts
 - the condition for pixel to be added to a region could be
 - neighbour pixels that are connected/adjacent due to similar values on average and/or variance, standard deviation

Region-based segmentation approaches (3)

- Split-and-Merge
 - it is a recursive method with an initial split step and a final merging step
 - the first step starts by dividing randomly the image into a random number of regions
 - then, each of those regions is recursively analysed and divided until homogeneous regions are obtained
 - according to some specified criteria
 - in the second step, contiguous regions, identified in the first step as being disjoint, are analysed and merged if together they form an homogeneous region
 - according to some specified criteria

References

- <https://courses.cs.washington.edu/courses/cse576/book/ch6.pdf>
- <http://www.uio.no/studier/emner/matnat/ifi/INF4300/h09/undervisningsmateriale/hough09.pdf>
- <https://web.fe.up.pt/~campilho/PDI/NOTES/EdgeDetection.pdf>
- <http://www.inf.ufg.br/~fabrizzio/mestrado/pdi/aulas/aula10.pdf>
- <http://gec.di.uminho.pt/lesi/vpc0304/Aula07Segmentação.pdf>
- http://iris.sel.eesc.usp.br/sel886/Aula_6.pdf
- <http://mtc-m18.sid.inpe.br/col/sid.inpe.br/mtc-m18%4080/2010/06.22.18.13/doc/106003.pdf?metadataarepository=&mirror=iconet.com.br/banon/2005/09.28.12.40>
- <https://www.mathworks.com/help/images/functionlist.html>